

### Task 3

The query was:

Select count(distinct (employees.emp\_no)) From Salaries,  
employees where employees.emp\_no = Salaries.emp\_no and  
Salaries > 50000 and Salaries < 70000 and to-date - from-  
date < 1


To add an index the query was:

Create INDEX ind-salary ON Salaries (Salary);

Table showing results of running the query with and without the index on Salaries

Number of Execution	Execution time (without index) in milliseconds	Execution time (with index) in milliseconds
1	1482	1426
2	1480	1395
3	1453	1397
4	1449	1398
5	1441	1395
6	1446	1397
7	1450	1396
8	1442	1402
9	1462	1393
10	1443	1390
12	1458	1406
13	1442	1399
14	1445	1396
15	1443	1394
16	1447	1396
17	1446	1394
18	1434	1392
19	1440	1395
20	1436	1394

A line graph Showing the execution times (in milliseconds) for the given query (when indexed and not indexed) for 20 successive runs whilst the query remains constant



Key	
—	Without index on Salaries
—	With index on Salaries

### Paragraph

When running the query without the index, most (if not all) execution times were above 1440 milliseconds. When running the query with an index, most execution times were around 1395. When running the query with an index on Salary, the ~~an~~ execution time was quicker for every single number of execution when compared to running without an



index. The highest time without an index is 1482 milliseconds.

The shortest time without an index is 1436 milliseconds.

The highest and shortest times with an index were 1426 and ~~1426~~ 1390 milliseconds respectively.

$1482 - 1390 = 92$  milliseconds. This is the time difference

from the highest to lowest execution times. From the results, we can conclude that the execution time is quicker when there is an index on Salaries compared to no index on Salaries.

We can also conclude that the execution time with an index stays more constant between executions as opposed to without an index on Salaries.

## Task 4


The query was:

Explain Select employees.first-name, ~~employees~~<sup>Salaries.</sup> Salary From  
employees INNER JOIN Salaries on employees.emp-no = Salaries.  
emp-no ORDER BY Salary ASC LIMIT 5;

~~The salaries table was queried~~

Rows	Employees	Salaries
Rows	299778	4

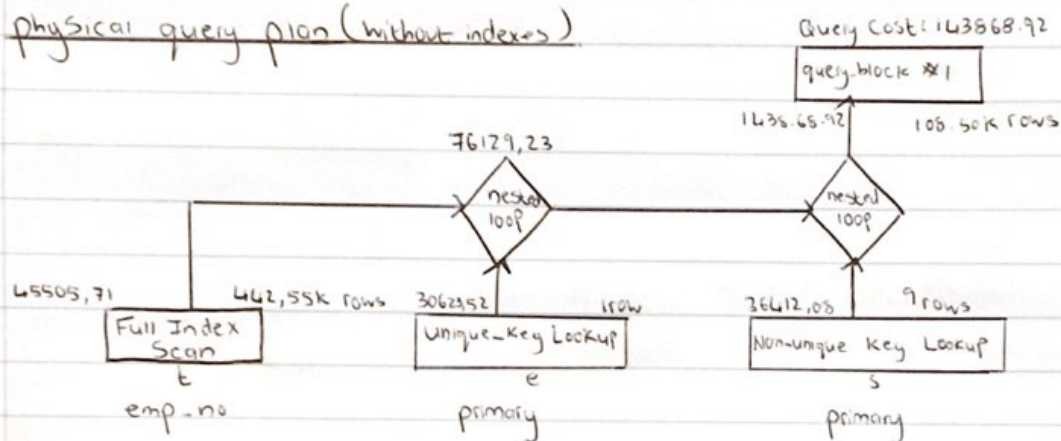
A histogram showing the number of rows queried for each table while the query remained constant



## Task 5

Query: Select e.first-name, e.last-name, t.title, s.salary  
from employees e, titles t, salaries s, where e.emp-no =  
t.emp-no ~~AND~~ and t.emp-no = s.emp-no and t.title =  
'engineer' and s.salary > 80 000 and s.salary < 100 000.


### physical query plan (without indexes)



Two indexes you can add could be on the titles and on employees

- Create Index employeeIdx on employees (emp-no);
- Create Index titleIdx on titles (title);

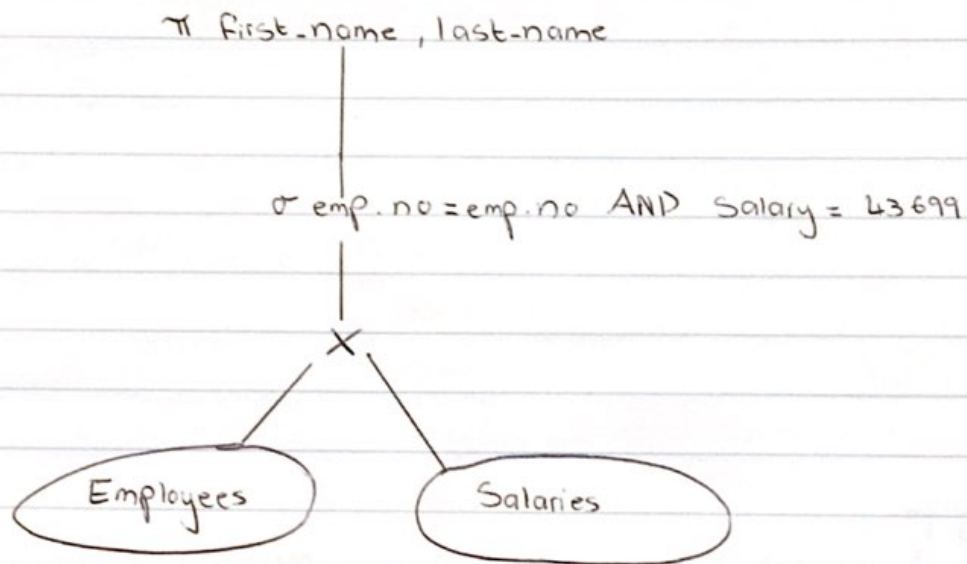
### physical query plan (with indexes)




### Task 6

Select distinct (first-name), last-name from employees, Salaries  
Where employees.emp-no = Salaries.emp-no and Salary > 43699;

### Query tree



### Optimised query tree





### Relational algebra Statement

$\pi$  first-name, last-name (employees  $\bowtie$  Salaries  $\cdot$  emp-no = employees  $\cdot$  emp-no  
( $\sigma$  salary = 43699 (Salaries)))

### Paragraph

Yes, I expected the optimised query tree execution times to be slightly better than the initial query tree execution times. I expected the time for optimisation queries to be much faster though. But nevertheless, faster times were achieved in the end.